

Structured Data with LLMs Done Right

A Practical Guide to Text Classification, Information Retrieval and Generation with LLMs

Jesus Villota

CEMFI

jesus.villota@cemfi.edu.es

November 11, 2025

Outline

Introduction

The Problem with Current Approaches

The Solution: Function Calling Schemas

Task 1: Text Classification

Task 2: Information Retrieval

Task 3: Structured Generation

Implementation & Best Practices

Conclusion

Three Core Tasks in Text Analysis

There are three fundamental ways of mapping text to structured data:

1. Text Classification

Assign predefined labels from a finite set of categories to text units.

$$f : \mathcal{D} \rightarrow \mathcal{C}$$

Example: Classify central bank statements as hawkish, dovish, or neutral

2. Information Retrieval

Extract specific data fields or entities from text.

$$g : \mathcal{D} \rightarrow \mathcal{V}_1 \times \mathcal{V}_2 \times \cdots \times \mathcal{V}_m$$

Example: Extract GDP forecasts, policy rates, and dates from Fed minutes

3. Structured Generation

Create new textual content adhering to a predefined schema.

$$h : \mathcal{D} \times \mathcal{S} \rightarrow \mathcal{T}$$

Example: Generate policy briefs with standardized sections

Motivation

LLMs are transforming empirical economics research:

- Central bank communication analysis (monetary policy stance)
- Financial sentiment from news and earnings calls
- ESG classification of corporate disclosures
- Contract information extraction (loan agreements, M&A)
- Policy text processing and analysis
- Literature review automation

The Problem

Current approaches relying on natural language prompting are fundamentally unreliable
— threatening research validity and reproducibility

The Problem with Current Approaches

The Standard Approach: Prompting

Most researchers use natural language prompting:

```
prompt = f"""
Classify the following text as 'positive', 'negative', or 'neutral'.
Return the result in JSON format:
{{
    "sentiment": "positive",
    "confidence": 0.85
}}

Text: {text}
"""

response = llm.generate(prompt)
```

Why This Fails

The model is **instructed** but not **constrained** to follow the format.

Reliability Crisis


LLMs are stochastic processes ⇒ they often fail to follow instructions:

- ✗ Return text that is not valid JSON
- ✗ Use different key names than specified
- ✗ Omit required fields or add unexpected fields
- ✗ Return values outside predefined categories
- ✗ Produce inconsistent formats across different inputs

Impact on Research

- ✗ **Broken** downstream processing **pipelines**
- ✗ Manual intervention introduces **bias**
- ✗ Poor **reproducibility** across model versions
- ✗ Impossible to **compare models** fairly

The Reliability Gap: Empirical Evidence



- **Prompting alone:** 85.1% adherence to schema instructions (GPT-4o)
- **Enriching prompt with structured schema:** 93% adherence to schema instructions (GPT-4o)
- **Structured schemas with strict mode:** 100% adherence (GPT-4o)

The Solution: Function Calling Schemas

What are Function Calling Schemas?

Function calling (also: "tool use" or "tool choice") allows researchers to define precise schemas that **guarantee** structured, parseable results.

Brief History

- Pioneered by OpenAI in 2023 with Chat Completions API
- Quickly adopted by all major LLM providers:
 - OpenAI, Anthropic Claude, Google Gemini, xAI Grok
 - Mistral, Cohere, GroqCloud, Azure OpenAI
- **Strict mode** achieves 100% schema adherence
- Now an industry standard for structured output

Transform structured output from a **probabilistic formatting challenge** into a **deterministic structural contract**

Unified Framework Structure

Any structured output task can be formulated as a function calling (JSON) schema:

```
tools = [{}  
  "type": "function",  
  "function": {  
    "name": "function_name",           # What this function does  
    "description": "Brief description",  
    "strict": true,                   # Enforce 100% adherence  
    "parameters": {  
      "type": "object",  
      "properties": {  
        "property_1": {  
          "type": "string",           # string, number, boolean, etc.  
          "description": "Instructions for property_1"  
        },  
      },  
      "required": ["property_1"],  
      "additionalProperties": false  
    }  
  }  
]
```

Key Schema Components

Core Elements

- **name**: Function name
- **description**: What it does
- **properties**: Output fields
- **required**: Mandatory fields

Available Types

- **string**: Text, dates, names
- **number**: Continuous values
- **integer**: Discrete counts
- **boolean**: Yes/no flags
- **array**: Lists of items
- **object**: Nested structures


Field Properties

- **type**: Data type (string, number, boolean, array, object)
- **description**: Instructions for this field
- **enum**: Restrict to specific values

Critical Setting

- **strict**: Enforce adherence
- **additionalProperties: false**
Prevents unexpected fields

How It Works: API Flow



Separation of Concerns

- **Schema defines** what information to return and in what format
- **Model determines** the actual content based on text analysis
- **Strict mode guarantees** format compliance

Task 1: Text Classification

Classification Schema Template

Key element: **enum** constraint restricts model to predefined categories

```
tools = [{
    "type": "function",
    "function": {
        "name": "classify_<task>",
        "description": "Classify text according to <criterion>",
        "strict": true,
        "parameters": {
            "type": "object",
            "properties": {
                "classification": {
                    "type": "string",
                    "enum": ["c_1", "c_2", "c_3", ..., "c_k"],
                    "description": "Detailed classification instructions"
                },
                "confidence": {"type": "number", "minimum": 0, "maximum": 1},
                "reasoning": {"type": "string"}
            },
            "required": ["classification", "confidence", "reasoning"],
            "additionalProperties": false
        }
    }
}]
```

Example: Classify FOMC statements by monetary policy stance

```
tools = [{
    "type": "function",
    "function": {
        "name": "classify_monetary_policy_stance",
        "description": "Classify central bank communication",
        "strict": true,
        "parameters": {
            "type": "object",
            "properties": {
                "policy_stance": {
                    "type": "string",
                    "enum": ["hawkish", "dovish", "neutral"],
                    "description": """Classify as:
                    - 'hawkish': preference for raising rates, inflation concerns
                    - 'dovish': preference for lowering rates, growth concerns
                    - 'neutral': balanced view or no clear direction"""
                },
                "confidence": {"type": "number", "minimum": 0, "maximum": 1},
                "reasoning": {"type": "string"}
            },
            "required": ["policy_stance", "confidence", "reasoning"],
            "additionalProperties": false
        }
    }
}]
```

Implementation

```
text = """Given the persistent upward pressure on prices  
and the need to ensure inflation expectations remain well-anchored,  
the Committee judges that a further tightening of monetary policy  
is warranted to bring inflation back to our 2% target."""  
  
response = client.chat.completions.create(  
    model="gpt-4o",  
    messages=[  
        {"role": "system", "content": "Expert in monetary policy."},  
        {"role": "user", "content": f"Classify: {text}"}  
    ],  
    tools=tools,  
    tool_choice={"type": "function",  
                "function": {"name": "classify_monetary_policy_stance"}},  
    temperature=0.0  
)  
  
result = json.loads(response.choices[0].message.tool_calls[0]  
                    .function.arguments)  
# Result: {"policy_stance": "hawkish", "confidence": 0.95, ...}
```

Guaranteed valid output every time!

More Classification Examples

Financial Sentiment

```
"sentiment": {  
    "type": "string",  
    "enum": ["positive",  
             "negative",  
             "neutral"]  
},  
"intensity": {  
    "type": "integer",  
    "enum": [1, 2, 3]  
}
```

Applications: earnings calls, news articles, analyst reports

ESG Classification

```
"primary_dimension": {  
    "type": "string",  
    "enum": ["environmental",  
             "social",  
             "governance",  
             "none"]  
},  
"commitment_level": {  
    "enum": ["aspirational",  
             "committed",  
             "implemented",  
             "reported"]  
}
```

Applications: corporate disclosures, annual reports

Task 2: Information Retrieval

Information Retrieval Schema Template

```
tools = [{}  
    "type": "function",  
    "function": {  
        "name": "extract_data",  
        "description": "Extract specific information from text",  
        "strict": true,  
        "parameters": {  
            "type": "object",  
            "properties": {  
                "entity_name": {"type": "string"},  
                "date": {"type": "string",  
                    "description": "YYYY-MM-DD format. Use null if not found."},  
                "amount": {"type": "number"},  
                "confidence": {"type": "number", "minimum": 0, "maximum": 1}  
            },  
            "required": ["entity_name", "date", "amount"]  
            "additionalProperties": false  
        }  
    }  
]
```

Example: Extract forecasts and policy decisions from FOMC minutes

```
tools = [{
  "type": "function",
  "function": {
    "name": "extract_economic_indicators",
    "description": "Extract economic forecasts and policy decisions",
    "strict": true,
    "parameters": {
      "type": "object",
      "properties": {
        "gdp_forecast": {"type": "number"},
        "inflation_forecast": {"type": "number"},
        "unemployment_forecast": {"type": "number"},
        "policy_rate_change": {
          "type": "number",
          "description": "Rate change in basis points (e.g., 25 for 0.25%)"
        },
        "forecast_horizon": {"type": "string"},
        "confidence": {"type": "number", "minimum": 0, "maximum": 1}
      },
      "required": ["policy_rate_change", "forecast_horizon", "confidence"],
      "additionalProperties": false
    }
  }
}]
```

Implementation

```
text = """The Committee decided to raise the target range for  
the federal funds rate by 25 basis points to 5.25-5.50 percent.  
Participants project GDP growth of 2.1% in 2024 and inflation  
to decline to 2.6% by year-end."""  
  
response = client.chat.completions.create(  
    model="gpt-4o",  
    messages=[  
        {"role": "system", "content": "Extract economic indicators."},  
        {"role": "user", "content": text}  
    ],  
    tools=tools,  
    tool_choice={"type": "function",  
                "function": {"name": "extract_economic_indicators"}},  
    temperature=0.0  
)  
  
result = json.loads(response.choices[0].message.tool_calls[0]  
                    .function.arguments)  
# Result: {"gdp_forecast": 2.1, "inflation_forecast": 2.6,  
#           "policy_rate_change": 25, "forecast_horizon": "2024", ...}
```

More Information Retrieval Examples

Loan Contract Terms

```
"loan_amount": {  
    "type": "number"  
},  
"interest_rate": {  
    "type": "number"  
},  
"maturity_date": {  
    "type": "string"  
},  
"collateral_type": {  
    "enum": ["real_estate",  
             "equipment",  
             "unsecured", ...]  
},  
"covenants_present": {  
    "type": "boolean"  
}
```

Corporate Events

```
"event_type": {  
    "enum": ["merger",  
            "acquisition",  
            "earnings",  
            "dividend", ...]  
},  
"event_date": {  
    "type": "string"  
},  
"companies_involved": {  
    "type": "array",  
    "items": {"type": "string"}  
},  
"transaction_value": {  
    "type": "number"  
}
```

Applications: Credit markets, event studies, financial databases

Task 3: Structured Generation

Structured Generation Schema Template

```
tools = [{}  
    "type": "function",  
    "function": {  
        "name": "generate_output",  
        "description": "Generate structured content",  
        "strict": true,  
        "parameters": {  
            "type": "object",  
            "properties": {  
                "title": {"type": "string", "description": "Concise title (max 15 words)"},  
                "summary": {"type": "string", "description": "Brief summary (max 200 words)"},  
                "key_points": {  
                    "type": "array",  
                    "items": {"type": "string"},  
                    "description": "List of 3-5 key points",  
                    "minItems": 3,  
                    "maxItems": 5  
                },  
            },  
            "required": ["title", "summary", "key_points"]  
        }  
    }  
}]
```

Example: Generate standardized policy briefs from legislation

```
tools = [{

    "type": "function",
    "function": {
        "name": "generate_policy_brief",
        "description": "Generate structured policy brief",
        "strict": true,
        "parameters": {
            "type": "object",
            "properties": {
                "title": {"type": "string", "description": "max 15 words"},
                "executive_summary": {"type": "string", "description": "max 150 words"},
                "affected_sectors": {"type": "array", "items": {"type": "string"}},
                "fiscal_impact": {
                    "type": "object",
                    "properties": {
                        "estimated_cost": {"type": "number"},
                        "revenue_impact": {"type": "number"},
                        "time_horizon": {"type": "string"}
                    },
                    "required": ["time_horizon"]
                },
                "key_provisions": {
                    "type": "array",
                    "items": {"type": "string"},
                    "minItems": 3,
                    "maxItems": 7
                }
            },
            "required": ["title", "executive_summary", "affected_sectors", "key_provisions"]
        }
    }
}]
```

Key Features for Generation

Length Constraints

Specify in descriptions: "max 200 words", "3-5 bullet points"

Array Constraints

Use `minItems` and `maxItems` to enforce list lengths

- Ensures consistent structure across generated outputs
- Prevents overly brief or verbose lists

Nested Objects

Use `type: "object"` for complex hierarchical structures

- Example: `fiscal_impact` contains multiple related fields
- Maintains logical grouping of related information

More Generation Examples

Company Profiles

```
"financial_highlights": {  
    "type": "object",  
    "properties": {  
        "revenue": {"type": "number"},  
        "net_income": {"type": "number"},  
        "year": {"type": "integer"}  
    }  
},  
"key_risks": {  
    "type": "array",  
    "items": {"type": "string"},  
    "minItems": 3,  
    "maxItems": 5  
}
```

Application: Cross-firm comparison

Literature Summaries

```
"main_research_question": {  
    "type": "string"  
},  
"key_findings": {  
    "type": "array",  
    "items": {"type": "string"},  
    "minItems": 3,  
    "maxItems": 5  
},  
"limitations": {  
    "type": "array",  
    "items": {"type": "string"},  
    "minItems": 2,  
    "maxItems": 4  
}
```

Application: Meta-analysis prep

Applications: Research automation, systematic reviews, database creation

Implementation & Best Practices

Development Workflow

1. Pilot Testing

- Test schema on 30-50 labeled examples
- Validate accuracy and refine descriptions
- Use reasoning field to understand decisions

2. Iterative Refinement

- Adjust category descriptions based on errors
- Clarify ambiguous cases in field descriptions
- Add examples in descriptions when helpful

3. Human Validation

- Validate random sample (10-20%) against human coding
- Assess reliability and identify systematic errors
- Document agreement metrics in paper

4. Full Deployment

- Process complete dataset
- Monitor confidence scores
- Flag low-confidence outputs for review

Model Selection and Configuration

Model Comparison

- Test multiple models on validation set
- Consider accuracy-cost tradeoffs
- Options: GPT-4o, Claude Sonnet, Gemini, Llama, etc.
- Use smaller models (GPT-4o-mini) for testing

Operational Considerations

- **Batch Processing:** Use batch APIs for cost savings
- **Rate Limiting:** Implement proper throttling
- **Error Handling:** Retry logic for API failures
- **Caching:** Cache common prompts when supported

Temperature Settings

- `temperature=0.0` for classification & extraction (max consistency)
- `temperature=0.3-0.5` for generation (some variation OK)

Cost Management

Estimate total costs before full implementation using token counts

Reproducibility Checklist

Document in Your Paper

- ✓ Exact model version (e.g., gpt-4o-2024-08-06)
- ✓ Complete schema with all properties and descriptions
- ✓ System prompts and user message templates
- ✓ All API parameters (temperature, tool_choice, etc.)
- ✓ Fixed seed parameter when available

Replication Materials

- ✓ Store complete schemas in supplementary materials
- ✓ Include example API calls
- ✓ Document any schema modifications during research
- ✓ Version output data if schema changes
- ✓ Report validation metrics (human agreement, confidence distributions)

Cross-Provider Compatibility

Function calling is supported by all major LLM providers:

- OpenAI (GPT-4, GPT-4o)
- Anthropic Claude (Sonnet, Opus)
- Google Gemini
- xAI Grok
- Mistral AI
- Cohere
- GroqCloud (fast inference)
- Azure OpenAI
- Local models (via APIs)

Portability

The JSON schemas in this paper work across providers with minimal syntax adjustments

- Core schema structure is universal
- API call syntax varies slightly
- Consult provider documentation for specifics

Conclusion

Key Takeaways

1. The Problem is Fundamental

- Prompting alone: 35-85% format adherence
- This is unacceptable for scientific research

2. The Solution: Function Calling Schemas

- Strict mode: 100% format adherence
- Industry standard across all major providers
- Deterministic formatting + analytical flexibility

3. Universal Framework

- Same structure for classification, retrieval, generation
- Easily combined for multi-task applications
- Ready-to-use templates for common economic research tasks

4. Reproducible Research

- Document model version, complete schema, parameters
- Output format is deterministic
- Fair model comparisons possible

Call to Action

Stop relying on prompting alone

Adopt function calling schemas

for reliable, reproducible, rigorous research

Thank You!

Questions?

Jesus Villota-Miranda

CEMFI

jesus.villota@cemfi.edu.es

Structured Data with LLMs Done Right

A Practical Guide to Text Classification, Information Retrieval and Generation